

The Platform and the Public Square: Social Media, Politics, Identity and Truth

How a handful of technology companies came to control the information environment — and what that means for democracy, identity, and the nature of truth itself.

filter bubble · echo chamber · misinformation · disinformation · algorithm · engagement · polarisation · radicalisation · surveillance capitalism · content moderation · confirmation bias · epistemic · democratic discourse · deepfake · bot · astroturfing · geopolitical · propaganda · tribalism · amplification

PART ONE · THE UNPRECEDENTED SCALE OF SOCIAL MEDIA

In the history of human communication, no technology has achieved the scale of modern social media platforms in so short a time. Facebook, founded in 2004, now has approximately three billion active users — more than a third of the entire population of the planet. YouTube, Instagram, TikTok, X (formerly Twitter), WhatsApp, and WeChat collectively reach billions more. For the first time in history, a significant proportion of humanity shares a common information environment — one that is not controlled by governments, not regulated by professional editorial standards, and not subject to the norms of accuracy and accountability that governed traditional media for most of the twentieth century. The implications of this shift are only beginning to be understood.

Social media platforms are, at their core, advertising businesses. Their revenue depends on keeping users on the platform for as long as possible, because longer engagement translates directly into more advertising impressions and therefore more revenue. This creates a fundamental tension at the heart of the social media model: the content that maximises engagement is not necessarily the content that is most accurate, most constructive, or most beneficial to individuals or to society. Research consistently shows that content triggering strong emotional reactions — particularly outrage, fear, and moral indignation — spreads faster and further than calm, measured, or nuanced content. The algorithm does not care about the truth. It cares about attention.

This is what the scholar Shoshana Zuboff has called surveillance capitalism — an economic system in which human experience and behaviour are harvested as raw data, processed into predictions about future behaviour, and sold to advertisers who wish to influence that behaviour. Every click, every pause, every like, every search is a data point that feeds the algorithm, which uses it to build an increasingly precise model of who you are, what you want, and — crucially — what kind of content will keep you scrolling. The result is a system that is extraordinarily effective at capturing attention, and that has reshaped the political and cultural landscape of the early twenty-first century in ways that were neither anticipated nor, for the most part, intended.

3 bn+

monthly active users on
Facebook alone

6 hrs

average daily time
online globally per
person

70%

of Americans get at
least some news from
social media

6×

more likely to share
false news than true
news

PART TWO · FILTER BUBBLES, ECHO CHAMBERS AND THE FRACTURING OF SHARED REALITY

One of the most important — and most debated — concepts in the study of social media's political effects is the **filter bubble**. The term, coined by internet activist Eli Pariser in 2011, describes the way in which personalisation algorithms curate each user's feed to show them content that reflects and reinforces their existing preferences, beliefs, and behaviours. Because the **algorithm** has learned what keeps you engaged, it serves you more of the same — more content from people you agree with, more stories that confirm your existing worldview, and progressively less exposure to perspectives that challenge, complicate, or contradict your existing beliefs.

The **filter bubble** feeds into what social psychologists call the **echo chamber** — an environment in which a particular set of beliefs is constantly repeated and reinforced, while alternative viewpoints are rarely encountered and, when they are, are typically encountered in a hostile or distorted form. In an echo chamber, **confirmation bias** — the natural human tendency to seek out and give greater weight to information that confirms what one already believes — is systematically amplified. Over time, the views held within the chamber tend to become more extreme, because the social rewards of expressing agreement are high and the social costs of expressing doubt or disagreement are significant.

The political consequences of this dynamic are severe. When different groups of people inhabit genuinely different information environments — seeing different news stories, encountering different facts, and receiving fundamentally different explanations of events — the capacity for **democratic discourse** is significantly undermined. Democracy, at its core, depends on a shared **epistemic** foundation: citizens need to agree, at least in broad terms, on what has actually happened and what the relevant facts are, before they can meaningfully disagree about what should be done. When that shared foundation fractures — when political opponents inhabit effectively different realities — compromise, negotiation, and democratic problem-solving become extraordinarily difficult.

"If I look at your Facebook page and you look at mine, we might as well be living on different planets. We have different facts, different stories, different heroes, different villains. That is not a recipe for democracy — it is a recipe for civil war." — Eli Pariser, author of The Filter Bubble

A landmark 2018 study published in the journal *Science* by researchers at MIT analysed 126,000 news stories shared on Twitter over a decade. The findings were striking: false news stories spread significantly faster, further, and more broadly than true ones. A false story was 70 per cent more likely to be retweeted than a true story. The primary driver was not bots or automated accounts — it was human beings, who found false stories more novel and emotionally engaging than accurate ones. The implications are deeply troubling: the architecture of social media appears to give **misinformation** a structural advantage over truth.

PART THREE · MISINFORMATION, DISINFORMATION AND THE WEAPONISATION OF TRUTH

It is important to distinguish between **misinformation** and **disinformation**, because they involve different actors and different degrees of intent. Misinformation refers to false or inaccurate information that is shared without the intention to deceive — the person spreading it genuinely believes it to be true. Disinformation, by contrast, is deliberately false information created and spread with the explicit intention of deceiving its audience — **propaganda** in the original sense of the word. Both are deeply problematic, but disinformation poses a qualitatively different threat because it involves organised, strategic manipulation of the information environment.

The tools available for creating and spreading **disinformation** have expanded dramatically in recent years. Deepfakes — highly realistic synthetic videos generated by artificial intelligence — can make public figures appear to say things they never said, with a level of visual authenticity that makes detection extremely difficult for the average viewer. Bots — automated accounts programmed to post, like, share, and interact at scale — can artificially amplify fringe views until they appear to represent mainstream opinion. **Astroturfing** — the creation of fake grassroots campaigns that appear to represent genuine popular movements — allows well-funded actors to manufacture the appearance of broad public support for positions that may actually be held by very few people.

The **geopolitical** dimensions of this problem became impossible to ignore following the 2016 US presidential election, when investigations revealed that Russian state-sponsored actors had conducted a sophisticated **disinformation** campaign on social media, creating thousands of fake accounts designed to inflame existing social divisions, suppress voter turnout among certain demographics, and amplify content favourable to particular political outcomes. Similar operations have since been documented in elections across Europe, Latin America, and Asia. The social media **platform** has become a terrain of geopolitical contestation — a space where state and non-state actors compete to shape public perception, often with goals that have nothing to do with the democratic process.

■ CASE STUDY — The 2016 US Election and Russian Disinformation

The Internet Research Agency, a Russian troll farm funded by the Kremlin, created approximately 3,500 Facebook advertisements seen by an estimated 126 million Americans. The campaign operated hundreds of fake accounts posing as American civic groups on both the left and the right, organising real-world events, inflaming racial and cultural divisions, and producing content specifically designed to reduce confidence in democratic institutions. Crucially, the operation did not aim primarily to support any particular candidate — it aimed to deepen polarisation and undermine the shared epistemic foundations on which democratic discourse depends. In this, by most assessments, it was highly effective.

PART FOUR · POLARISATION, RADICALISATION AND THE POLITICS OF OUTRAGE

The relationship between social media and political polarisation is one of the most actively researched questions in contemporary social science. The evidence is mixed and contested, but a number of robust findings have emerged. First, polarisation — the tendency for political attitudes to become more extreme and for partisan groups to regard each other with increasing hostility — has increased sharply in many democracies over the same period that social media has become dominant. While it would be a serious oversimplification to attribute this trend entirely to social media (economic inequality, cultural change, and the decline of traditional media all play significant roles), the evidence that social media accelerates and deepens polarisation is substantial.

The mechanism is partly a function of how the engagement algorithm rewards outrage. Content that provokes anger, contempt, or moral indignation generates more comments, shares, and interactions than content that provokes reflection or nuanced thought. As a result, political content on social media tends to drift towards the extreme — because extreme content is more emotionally engaging, and emotional engagement is what the algorithm rewards. Users who begin by engaging with mainstream political content are gradually served increasingly extreme material, a process that researchers have documented extensively on platforms including YouTube, where the recommendation algorithm has been shown to guide users towards progressively more radical content in a process that has become known as radicalisation by recommendation.

Perhaps the most profound effect of this dynamic is what it has done to political tribalism — the tendency to view politics not as a pragmatic process of problem-solving and compromise, but as an existential conflict between groups with fundamentally incompatible values and interests. When political opponents are consistently depicted — by the algorithm and by the partisan content it promotes — as not merely wrong but as evil, corrupt, and dangerous, the normal psychological preconditions for democratic compromise are eroded. In a polarised social media environment, the act of reaching across political lines becomes not a virtue but a betrayal — an act of collaboration with the enemy. This is, by any reasonable assessment, a serious threat to the health of democratic institutions.

■ THE CASE FOR SOCIAL MEDIA IN DEMOCRACY

- Social media has given marginalised voices a platform that traditional media denied them.
- Citizen journalism and viral amplification have exposed injustices that mainstream media ignored.
- Political movements — from the Arab Spring to #MeToo — were organised and sustained through social platforms.
- Information access has been democratised — anyone can now reach a global audience.
- Social media enables rapid coordination of civil society responses to crises.

■■ THE CASE AGAINST SOCIAL MEDIA IN DEMOCRACY

- Algorithmic amplification gives extreme voices disproportionate reach and influence.
- The speed of social media prevents the verification processes that accurate journalism requires.
- Authoritarian regimes have used the same platforms to suppress dissent and spread state propaganda.
- The democratisation of publishing has also democratised disinformation at unprecedented scale.
- Social media profit models are structurally misaligned with democratic health.

PART FIVE • IDENTITY, PERFORMANCE AND THE CURATED SELF

The effects of social media are not limited to the political sphere. At a more intimate and personal level, the platforms have fundamentally altered the way in which people construct, present, and experience their own identities. The sociologist Erving Goffman argued in the 1950s that all social interaction involves a form of performance — people consciously and unconsciously manage the impressions they make on others. Social media has taken this insight and industrialised it, creating environments in which the performance of identity is not merely a side effect of social interaction but its primary purpose.

The Instagram feed, the carefully composed LinkedIn profile, the Twitter persona — these are not simply records of a life as it is lived. They are curated constructions — edited, filtered, and optimised for a particular audience. The research on the psychological consequences of this is substantial and troubling. Studies consistently show that heavy social media use, particularly platforms centred on image and lifestyle, is associated with increased rates of anxiety, depression, body image issues, and loneliness — especially among young women. The mechanism appears to be social comparison: the curated highlight reels of others' lives create an implicit standard against which one's own, inevitably more ordinary life compares unfavourably.

The influencer economy has professionalised this dynamic, creating a new class of media figures whose income depends on the size and engagement of their following, and who therefore have powerful commercial incentives to present an idealised version of their lives. The authenticity that influencers trade on — their apparent willingness to share personal struggles, vulnerabilities, and behind-the-scenes moments — is itself a carefully managed performance, subject to the same editorial judgements and commercial calculations as any other form of media production. The paradox of the authentic influencer is that their authenticity is itself a product.

Social media companies occupy an extraordinarily uncomfortable position in the contemporary information landscape. They are, by any reasonable measure, publishers — they create, curate, and distribute content that reaches billions of people and has measurable effects on political outcomes, public health, and social cohesion. Yet they have consistently resisted the accountability frameworks that apply to traditional publishers, arguing that they are merely neutral platforms providing infrastructure for others' speech, rather than editorial actors responsible for the content they carry.

This position has become increasingly untenable. The decisions that platforms make about **content moderation** — what to remove, what to downrank, what to label as disputed or misleading — are editorial decisions with profound consequences. When Facebook chose to suppress the reach of news content in its feed, traffic to news publishers collapsed. When Twitter banned certain political accounts, it demonstrated that it could and would make judgements about what speech was acceptable on its **platform**. When YouTube demonetised certain categories of content, it effectively determined which kinds of speech could be economically viable. These are not the decisions of a neutral platform — they are the decisions of a powerful editorial actor that has so far resisted meaningful democratic oversight.

The question of how to regulate social media platforms — and who should do so — is one of the most contested policy debates of the current era. The European Union has moved furthest with its Digital Services Act, which imposes significant obligations on large platforms around transparency, **content moderation**, and algorithmic accountability. The United States, by contrast, has been largely paralysed by political disagreement and the extraordinary lobbying power of the technology industry. In the meantime, the platforms continue to wield **geopolitical** influence that dwarfs that of many nation states, largely without the accountability mechanisms that democratic governance normally demands.

◆ The Fundamental Tension — Speech vs Safety

Content moderation presents an inescapable dilemma: to remove harmful content aggressively is to risk censorship and the suppression of legitimate political speech; to moderate lightly is to allow **misinformation**, hate speech, and coordinated manipulation to flourish. There is no technical solution to this problem — it is fundamentally a question of values, and of who has the authority to make those value judgements. When a handful of private technology companies in Silicon Valley effectively determine the boundaries of acceptable speech for three billion people, the adequacy of current accountability frameworks is called seriously into question.

The social media platforms that emerged in the first two decades of the twenty-first century were built on a set of assumptions that have since been shown to be either naive or actively misleading: that more connectivity is always better; that the free flow of information is inherently democratic; that neutral **platform** technology is value-free; and that the aggregate of individual choices, however driven by

algorithmic nudges and engineered emotional triggers, constitutes a form of free expression. All of these assumptions have been challenged by the evidence of the past decade, and the debate about what to do about it is now one of the defining political and philosophical questions of our time.

Reform proposals range from the incremental to the radical. Some argue for enhanced transparency requirements — forcing platforms to disclose how their algorithms work, who funds political advertising, and what criteria govern **content moderation** decisions. Others advocate for structural changes to the business model itself, arguing that the advertising-driven **engagement** model is fundamentally incompatible with democratic health, and that platforms should be required to adopt subscription models or non-profit structures. A growing number of researchers and policymakers argue that the major platforms should be broken up under antitrust law, on the grounds that their current scale gives them a monopoly power over public discourse that no private company should possess.

What seems clear is that the current situation — in which the most powerful information distribution systems in human history are controlled by a small number of private companies, optimised for commercial **engagement** rather than democratic health, and subject to minimal effective accountability — is not sustainable. The question is not whether change will come, but whether it will be driven by democratic deliberation and public interest, or whether it will be imposed by the internal crises — of trust, of regulation, of political legitimacy — that the platforms' current trajectory seems likely to produce. The age of social media has barely begun, and its most consequential chapters may still be ahead of us.

■ TALKING POINTS

Do you think social media has been — on balance — good or bad for democracy? What evidence from the article supports your view?

The article distinguishes between misinformation and disinformation. Why does this distinction matter? Can you think of an example of each?

Should social media companies be regulated like publishers — with legal responsibility for the content they carry? What are the risks of this approach?

The article says the influencer's 'authenticity is itself a product'. Do you think audiences are aware of this — and does it matter if they are not?

Is it possible to use social media in a way that protects your epistemic independence — your ability to form views based on evidence rather than algorithmic curation?